

---

# The Fixed-Effect Fallacy and Claims about GPT-4 Capabilities

---

Thomas Ball<sup>1</sup> Shuo Chen<sup>1</sup> Cormac Herley<sup>1</sup>

## Abstract

Rapid improvements in the performance of large language models (LLMs) have spurred great interest in evaluating their capabilities. An accumulation of observed successes and failures at particular tasks unfortunately does little to settle questions about LLM reliability. We present measurements of GPT-4 performance on several deterministic tasks. We examine several conditions per-task and perform enough trials so that statistically significant differences can be detected. We find that seemingly trivial modifications in the task-prompt can yield differences far larger than can be explained by sampling effects. We conclude that efforts to quantify LLM capabilities easily succumb to the language-as-fixed-effect fallacy, where experimental observations are improperly generalized beyond what the data supports. A consequence appears to be that intuitions that have been formed based on interactions with humans form a very unreliable guide as to which input modifications should “make no difference” to LLM performance.

## 1. Introduction

Rapid improvements in the performance of large language models (LLMs) have spurred great interest in evaluating their capabilities. In addition to answering general knowledge questions and summarizing text, GPT-4 has demonstrated the capability to compose poetry, solve chess puzzles and Geometry problems, and perform basic coding tasks. Capabilities that seem beyond the simple next-token-prediction they were trained on, causes some to suggest this as evidence of emergent behaviors from LLMs, or even that we may be witnessing the early signs of Artificial General Intelligence (AGI) (Bubeck et al., 2023). Others are not convinced, and suggest that LLMs simply parrot pastiches of text snippets from their training sets (Bender et al., 2021).

---

<sup>1</sup>Microsoft Research, Redmond, WA. Correspondence to: Firstname1 Lastname1 <first1.last1@xxx.edu>, Firstname2 Lastname2 <first2.last2@www.uk>.

The documentation of surprising capabilities has been accompanied by many accounts of failures. Hallucinations (where LLMs offer plausible but entirely invented detail) have proved hard to eliminate. Arkoudas points out that GPT-4 struggles with some basic tasks that humans find easy or trivial; e.g., they aren’t reliable even on tasks such as counting, multiplication, etc (Arkoudas, 2023). McCoy et al suggest that many of the remarkable capabilities are simply artifacts of the training set and autoregressive task that GPT-4 was trained to solve (McCoy et al., 2023).

An accumulation of observed successes and failures at particular tasks unfortunately does little to settle questions about LLM reliability. In this paper we present results on a series of basic algorithmic tasks that all have deterministic answers, allowing us to assess performance without costly and subjective hand-labelling or assessments. They also involve inputs that can be sampled from a large population (e.g., lists of numbers, name-value pairs, or  $k$ -digit numbers). This allows us to automatically assess performance over large samples. It also enables us to modify the task, e.g., by rewording the prompt, to check which factors affect performance. For example, our population might be length-21 lists of floating point numbers, and the task might be to find the median, but modifications might be to try reworded versions of the prompt, or try lists with a different number of significant decimal places given.

Our contributions are as follows. We present measurements of GPT-4 performance on several deterministic tasks. We examine several conditions per-task and perform enough trials (500 per condition unless otherwise stated) so that statistically significant differences can be detected. For all tasks and all conditions this entails about 37k responses from GPT-4; all prompts, responses and associated meta-data are openly available to those who wish to check or build upon our findings.<sup>1</sup> We measure performance on tasks such as counting, sorting, multiplication, etc, and find that accuracy, while better-than-random, is often poor. We find that seemingly trivial modifications in the task-prompt can yield differences far larger than can be explained by sampling effects.

We conclude that efforts to quantify LLM capabilities easily succumb to the language-as-fixed-effect fallacy (Clark,

---

<sup>1</sup>Withheld for anonymized submission

1973; Coleman, 1964), where experimental observations are improperly generalized beyond what the data supports. A consequence appears to be that intuitions that have been formed based on interactions with humans, form a very unreliable guide as to which modifications should “make no difference” to LLM performance.

We wish to be clear that our goal is not to determine whether LLMs can or cannot count, sort, or multiply, etc. First, we have other ways of performing these tasks; second, it is possible that prompt engineering, the use of Chain-of-Thought reasoning, or the invocation of plug-ins might improve performance. Our goal, rather, is to draw attention to the fact that LLMs seem particularly susceptible to a major pitfall that exists when we go from particular experimental observations to general claims about capabilities. That is, sensitivity to seemingly trivial modifications means that observed accuracy numbers cannot be assumed to generalize (even to entirely equivalent versions of a task). This means that reporting observed performance or accuracy numbers on deterministic tasks (such as standardized tests (Katz et al., 2023; Takagi et al., 2023; Nori et al., 2023), textbook problems, etc) is not sufficient to establish a general capability.

## 2. Background: The Language-as-Fixed-Effect Fallacy

The Language-as-Fixed-Effect Fallacy, as described by Clark (Clark, 1973), is the phenomenon where a claim supported by statistical evidence does not generalize beyond the specifics of the experimental setup. He illustrates with a language-task thought-experiment originally proposed by Coleman (Coleman, 1964). Let  $N$  be the set of all English nouns,  $V$  the set of all verbs, and let  $T(\cdot)$  be a test statistic representing how well humans perform at some task involving words (e.g., how well they can spell them, how quickly they can type them, etc). Suppose that experimenter A wishes to test the hypothesis that people perform the task better on nouns than on verbs:

$$\mathcal{H}_A = T(N) > T(V).$$

Suppose experimenter B wishes to test the opposite:

$$\mathcal{H}_B = T(N) < T(V).$$

Let’s stipulate, by contrast, that they are both wrong, and that  $T(N) = T(V)$ .

As a test of  $\mathcal{H}_A$  the first experimenter selects subsets  $N_A \subset N$  and  $V_A \subset V$  each with some fixed number of randomly selected nouns and verbs. With this choice she recruits participants and on finding that  $T(N_A) > T(V_A)$ , by a statistically significant amount, she rejects the null hypothesis (that there’s no difference) and concludes she has firm evidence in favor of  $\mathcal{H}_A$ . Similarly, the second experimenter

selects at random different subsets  $N_B \subset N$  and  $V_B \subset V$  with fixed numbers of nouns and verbs. With this fixed choice he recruits participants and finds  $T(N_B) < T(V_B)$ , by a statistically significant amount, and concludes this is firm evidence for  $\mathcal{H}_B$ .

The problem is that while both A and B intend to generalize to the whole population  $N$  and  $V$  they have tested only on particular subsets. There is good evidence to believe that, with any collection of participants, we could verify both  $T(N_A) > T(V_A)$  and  $T(N_B) < T(V_B)$ , but neither of these is enough to support either  $\mathcal{H}_A$  or  $\mathcal{H}_B$ . In the language of statistical testing our experimenters have treated random effects as fixed (Clark, 1973).

Fixed effects are those that are considered constant across the relevant population, while random effects are those that vary (for an account of various other definitions see (Gelman, 2005)). In the experiments above there are two populations involved: the populations of noun-verb collections, and the population of human participants. When she generalized from  $N_A, V_A$  to  $N, V$  our first experimenter implicitly assumed that any other subsets  $N_C \subset N$  and  $V_C \subset V$  would also give the result that she observed (i.e.,  $T(N_C) > T(V_C)$ ). If this were true she’d be justified in thinking that her observed difference was powerful evidence for  $\mathcal{H}_A$ . If this is not true then her experiment supports only the narrow uninteresting claim  $T(N_A) > T(V_A)$ . Effectively, she assumed that what she observed wasn’t particular to  $N_A, V_A$  but general to  $N, V$ .

In a colloquial sense fixed effects are ones where the particular choice doesn’t affect the generality we wish to claim. We expect, for example, that what an experimenter had for breakfast or what color socks she was wearing has no effect on the outcome; these are not details that have to be faithfully reproduced to ensure replication of the original experiment. In this telling the fixed-effect fallacy is simply assuming that certain details don’t matter when in fact they do. Unfortunately, there’s no simple way to determine that a certain variable has no influence on an experimental result; experiments necessarily involve many judgements about which details matter and which do not, and many of those judgements are subjective. One of our findings is that intuitions about which modifications might make a difference can be very flawed; that human performance remains constant under a certain modification is no guarantee at all that LLM performance also will.

## 3. The Fixed-Effect Fallacy and LLM Task Performance

We wish to evaluate whether, and how well, an LLM can perform a particular task that has a single deterministic correct answer (e.g., counting). One approach might be to

produce a list of objects and prompt the LLM to count the occurrences of a particular item. To make the experimental setup concrete we might specify a list length and dictionary of possible elements. For example:

```
rLen = 20
listOfItems = ['mango', 'peach']
r = random.choices(listOfItems, k = rLen)
```

is a Python snippet that will return a length-20 list with the elements of `listOfItems` chosen at random with replacement. When there are only two elements, as shown, there’s a population of  $2^{20}$  such lists; call this population  $\mathbf{R}$ . We might prompt the LLM with:

```
prompt = ``How many times does `mango`
        appear in this list: `` + str(r)
```

where  $r \in \mathbf{R}$ . By repeating this query with many different elements of  $\mathbf{R}$  we might try to build a picture of the LLM’s performance at the task.

In this setup choice of list from  $\mathbf{R}$  is being treated as the only random effect; i.e., the only source of variation (Gelman & Hill, 2006). We are testing how well the LLM does over many different members of  $\mathbf{R}$  but are assuming that other factors we might vary make no difference. However, there are many other populations of lists that we might try, and there are many other wordings of the prompt that could be used. If we use observed success with the above prompt to conclude that our LLM can count elements of a length-20 list with a particular success rate we are implicitly assuming that these other possible choices would make no difference. For example, an alternative to the prompt above might be:

```
prompt = ``Here is a list: `` + str(r)
        + ``. How many times does `mango`
        appear on it?``
```

This would appear to be an equivalent evaluation of the task, or a modification that should make no difference. Unfortunately, this is not the case.

As we show in Section 4.2 these assumptions most definitely do not hold. Wording of the prompt and choice of the particular items to be counted can make a substantial difference to the answer (see Table 1). For example, the hypothesis that tests using the two prompts given above (with everything else held constant) produce results drawn from the same distribution, is robustly rejected by a  $\chi^2$  test. Thus, if we report that our LLM can count with a particular success rate we are committing the same fixed-effect error as experimenters A and B above.

When we encounter a particular experimental result (e.g.,  $q = 0.86$  (86.0%) on the  $N = 10$  counting task in Table 1) we generally understand that this involves some margin

of error. For example, rather than  $q$ , we expect a repeat of the experiment to produce an estimate  $q \pm \delta_q$ . A very familiar case exploits the fact that 95% of the values of a normal distribution lie within 1.96 standard deviations of the mean, so we can write  $\delta_q = 1.96 \cdot \sqrt{q \cdot (1 - q) / N}$  and be confident that 95% of trials will fall in this interval (Taylor & Thompson, 1982).

However, it is important to keep track of the baked-in assumptions: this estimate assumes that variance from sampling the list population (i.e., sampling  $\mathbf{R}$ ) is the only source of randomness. If significant other sources of randomness exist, then we know only that  $\delta_q$  is greater than (and possibly much greater than)  $1.96 \cdot \sqrt{q \cdot (1 - q) / N}$ . That is, we have only a lower bound on our margin-of-error. We can’t rule out, for example, that the 95% confidence interval is  $\pm 30\%$ . The results of Section 4 show that other sources of randomness that are too large to ignore do exist for several of the tasks we consider (and in some cases appear far greater than the variance due to sampling).

## 4. Tasks

### 4.1. Experimental setup

In order to test LLM performance we choose tasks that have deterministic answers, and where it is relatively easy to decide if the LLM gives the correct answer. This obviates the need for subjective assessments, heuristics, hand-labelling or error-prone parsing of the response, and allows us to scale-up testing. The tasks we examine are: counting, finding the maximum, median and sorted version of a list of numbers, long multiplication, and basic composite tasks combining the above tasks. The difficulty with counting and long multiplication has been observed by others (Arkoudas, 2023).

Unless otherwise specified all of the conditions were evaluated on 500 independent runs. Thus, for example, if a table entry reports a success rate of 89.0% on a task, and sampling were the only source of randomness, then a reasonable estimate of the 95% confidence interval would be  $1.96 \cdot \sqrt{0.89 \times 0.11 / 500} \approx 2.74\%$ . However, an important finding, below, is that there are significant other sources of randomness, and the conventional way of estimating margins-of-error cannot be applied. All of the trials are performed using the OpenAI GPT-4 API. The results of all queries are available in the GitHub repository.

For all of the tasks we give an example prompt together with the correct answer and GPT-4’s answer. In all of the examples shown the GPT-4 response is incorrect. This is not reflective of its accuracy: in each case we give a table showing how accuracy evolves with problem size. However, in giving examples where the answers are incorrect we illustrate that they are often very significantly better-than-

random.

## 4.2. Count

First we examine the capability of GPT-4 to perform basic counting tasks. We choose a length- $N$  list with two possible elements and ask GPT-4 to count the number of occurrences of the first element. An example query is (let's call this wording #1):

**Prompt:** How many times does 'mango' appear in this list: [mango, peach, peach, peach, mango, mango, mango, peach, peach, peach, mango, mango, mango].  
**Correct Answer:** 7  
**GPT-4 Answer:** 'Mango' appears 6 times in this list.

We evaluate for five different target lengths; the results are shown in the first column of Table 1. In choosing modifications of this task we choose two different variations of the input list by replacing the word-pair 'mango/peach' with 'airedale/aspidistra' (results in column 2) and the random strings 'xhsfgre/jnosdfr' (results in column 3). We also examine one simple rewording of the prompt (let's call this wording #2):

```
prompt = ``Here is a list: `` + str(r)
        + ``. How many times does 'mango'
        appear on it?``
```

This gives us a total of four conditions, all of which involve the same basic counting task. We evaluate each condition with list lengths  $N = 10, 15, 20, 30$  and  $40$ , and we perform 500 trials per condition. The results are shown in Table 1. Thus, the five rows and first four columns represent a total of  $5 \times 4 \times 500 = 10,000$  queries to GPT-4.

We use a  $\chi^2$  test to determine if the responses to different ways of phrasing the task are drawn from the same distribution. For example, we can take the null hypothesis to be that some row of the first and fourth columns of Table 1 represent answers drawn from the same distribution. E.g., for  $N = 10$  there were 445/500 and 483/500 correct trials respectively. Using a standard  $\chi^2$  goodness-of-fit to compare these two distributions of correct/incorrect answers yields ( $\chi^2 = 21.6$ ,  $df = 1$ ,  $p = 3.34e - 6$ ). The  $p$ -value can be taken as an estimate of the probability of these results being observed if columns 1 and 4 of row 1 were produced by the same process; generally when  $p < 0.05$  we say that the null hypothesis is rejected. Similarly for all the other rows, the hypothesis (that results of the task with different wording are drawn from the same distribution) is rejected. The degrees-of-freedom is  $df = 1$  for all of our tests since we are always doing pairwise comparisons on tasks on a binary outcome (Taylor & Thompson, 1982).

The results of our  $\chi^2$  tests are given in the right-hand side

of Table 1. The null hypothesis is robustly rejected for all lengths when comparing columns 1 and 4 (i.e., simply switching between wording #1 and wording #2 with the 'mango/peach' word-pair). The null hypothesis is rejected for several lengths when comparing columns 1 and columns 2, 3 (i.e., simply switching the word-pair while using wording #1). This demonstrates that simple modifications of the task (that might easily be assumed to make no difference) in fact are sources of variance beyond what can be explained by sampling effects.

We note also that the GPT answers are biased toward under-counting. For example in the 'mango/peach' case the mean of the correct answers for the five lengths tested (i.e.,  $N = 10, 15, 20, 30$  and  $40$ ) were: (5.57, 7.96, 10.57, 15.46, 20.6) and the GPT-4 answers were (5.45, 7.57, 10.04, 14.09, 18.5). Thus, across 500 trials, the mean GPT-4 answers were always lower. Among the 500 trials the ratio of over-counts:under-counts was (55 : 0, 192 : 2, 248 : 11, 428 : 10, 451 : 11).

## 4.3. Maximum, Median and Sort

Here we ask GPT-4 to perform elementary tasks on lists of numbers: return the maximum, median and sorted version of the list. As in Section 4.2, we evaluate four different conditions. First we ask for the maximum (or median or sorted version) of a list of  $N$  numbers drawn uniformly-at-random from the interval (100.0, 20000.00) and rounded to two decimals places. An example of the prompt for the median-finding task is:

**Prompt:** What is the median value in this list: [7176.36, 5222.86, 1089.62, 19927.36, 5655.72, 18355.58, 18978.7, 7028.49, 14190.57, 14243.69, 11251.69]. Please write 'Answer='  
**Correct Answer:** 11251.69  
**GPT-4 Answer:** 7176.36

Second, we repeat with the numbers rounded to 12 decimal places. Third, we repeat with integers drawn uniformly-at-random from (10, 99) (i.e., all list-members are 2-digit numbers). Finally, we use a list of  $N$  name-value pairs, where a randomly-chosen name is associated with a number drawn uniformly-at-random from the interval (100.0, 20,000.00) and rounded to two decimals places. An example of the latter query is:

**Prompt:** Please sort this list in ascending order: [John: \$12158.21, Mary: \$1416.51, Peter: \$7507.58, Vivek: \$10941.54, Xian: \$10530.84, Alex: \$1641.14, Maria: \$1025.49, Frank: \$260.85, Luis: \$7464.35, Manuel: \$1782.86, Kristen: \$10085.24].  
**Correct Answer:** [Frank: \$260.85, Maria: \$1025.49, Mary: \$1416.51, Alex: \$1641.14, Manuel:



$N$	mango/ peach	airedale/ aspidistra	xhsfgr/ jnosdfi	Prompt reworded	Compare Cols(1,2) ( $\chi^2, p$ )	Compare Cols(1,3) ( $\chi^2, p$ )	Compare Cols(1,4) ( $\chi^2, p$ )
10	89.0%	91.2%	87.8%	96.6%	(1.36, 2.44e-1)	(0.35, 5.54e-1)	<b>(21.61, 3.34e-6)</b>
15	61.2%	53.6%	54.8%	88.6%	<b>(5.91, 1.51e-2)</b>	<b>(4.2, 4.03e-2)</b>	<b>(99.84, 1.66e-23)</b>
20	48.2%	29.6%	42.6%	76.2%	<b>(36.39, 1.62e-9)</b>	(3.16, 7.53e-2)	<b>(83.36, 6.83e-20)</b>
30	12.4%	7.4%	21.4%	43.6%	<b>(7.01, 8.12e-3)</b>	<b>(14.42, 1.46e-4)</b>	<b>(120.71, 4.41e-28)</b>
40	10.2%	7.6%	14.2%	21.0%	(2.08, 1.49e-1)	(3.73, 5.33e-2)	<b>(22.15, 2.53e-6)</b>

Table 1. Percent correct for counting the occurrences of a length- $N$  list with two items chosen uniformly-at-random. Performance decays rapidly with list length. The first three columns use (prompt) wording #1 with lists formed using the word-pairs shown; the fourth column uses wording #2 and lists formed using the word-pair ‘mango/peach.’ On the right-hand side of the table we present  $\chi^2$  tests comparing the results of the first condition with each of the others. This test evaluates the null hypothesis that the answers in the various conditions are drawn from the same distribution. Boldface entries are cases where  $p < 0.05$  and we reject the null hypothesis. The null hypothesis is robustly rejected for all lengths when comparing columns 1 and 4 (i.e., simply switching between wording #1 and wording #2 with the ‘mango/peach’ word-pair). The null hypothesis is rejected for several lengths when comparing columns 1 and 2, 3 (i.e., simply switching the word-pair while using wording #1). This demonstrates that simple modifications of the task (that might easily be assumed to make no difference) in fact are sources of variance beyond what can be explained by sampling effects.

\$1782.86, Luis: \$7464.35, Peter: \$7507.58, Kristen: \$10085.24, Xian: \$10530.84, Vivek: \$10941.54, John: \$12158.21]

**GPT-4 Answer:** [Frank: \$260.85, Maria: \$1025.49, Mary: \$1416.51, Alex: \$1641.14, Manuel: \$1782.86, Peter: \$7507.58, Luis: \$7464.35, Kristen: \$10085.24, Xian: \$10530.84, Vivek: \$10941.54, John: \$12158.21]

The results of the maximum, median and sorting tasks are given in Tables 2, 3 and 4 respectively. The four different list conditions are explored in columns 1-4 of these tables. As in Section 4.2, we use a  $\chi^2$  test to explore whether these different variations on the task produce answers that appear drawn from the same distribution. The right-hand portion of Tables 2, 3 and 4 gives the results; we do goodness-of-fit tests to compare columns 2, 3 and 4 with column 1.

Table 2 shows the results of the maximum-finding task. Performance in all conditions is good, though not perfect (e.g, results in the first four columns are  $> 95.0\%$ ). The  $\chi^2$  tests show that the hypothesis that performance on the name-value version of the list is consistent with performance on the value-only list is rejected for all lengths. The hypothesis that performance on the integer version of the list is consistent with performance on the 2-decimal floats list is rejected for lengths 11 and 15.

Table 3 shows the results of the median-finding task. Performance in all conditions is poor (e.g, results in the first four columns are  $< 90.0\%$ ). The  $\chi^2$  tests show that the hypothesis that performance when the numbers are drawn from (10.0, 20000.0) is consistent with performance when numbers are drawn from (10, 99) is rejected for all lengths. The hypothesis that that name-value version of the list is consistent with performance on the value-only list is also re-

jected for all lengths. Note that the  $p$ -values in both cases are  $\ll 0.05$ , so the probability that the same process accounts for both conditions is very low.

Table 4 shows the results of the sorting task. Performance in conditions 1-3 is good, but is very poor in condition 4 (e.g, results in column 4 are  $< 65.0\%$ ). The  $\chi^2$  tests show that the hypothesis that performance when the numbers are drawn from (10.0, 20000.0) is consistent with performance when numbers are drawn from (10, 99) is rejected for all lengths. The hypothesis that that name-value version of the list is consistent with performance on the value-only list is also rejected for all lengths. Again, the  $p$ -values indicate robust rejection of these hypotheses.

#### 4.4. Long Multiply

Here we evaluate performance at long multiplication, where we prompt the LLM to calculate the product of a  $k_1$ -digit by a  $k_2$ -digit number. An example for  $4 \times 4$  is:

**Prompt:** What is the product of 6438 and 9038?  
Please write ‘Answer =’  
**Correct Answer:** 58186644  
**GPT-4 Answer:** Answer = 58169844.

Table 5 shows the performance multiplying a  $k_1$ -digit by a  $k_2$ -digit number for  $k_1, k_2 \in \{2, 3, 4, 5\}$ . Apart from the  $2 \times 2$  case the results are largely poor. Observe that perfect performance on the  $2 \times 2$  task drops to negligibly correct answers for  $4 \times 4$ .

Since there is sometimes a significant difference between the  $k_1 \times k_2$  result with the  $k_2 \times k_1$  result we perform a  $\chi^2$  test on several of the off-diagonal elements. The results are shown in Table 6. Note that results for the  $4 \times 2$  and  $2 \times 4$  are significantly different, as are those for  $5 \times 2$  and  $2 \times 5$ .

## The Fixed-Effect Fallacy and Claims about GPT-4 Capabilities

	Col. 1	Col. 2	Col. 3	Col. 4	Compare Cols(1,2) ( $\chi^2, p$ )	Compare Cols(1,3) ( $\chi^2, p$ )	Compare Cols(1,4) ( $\chi^2, p$ )
11	98.0%	99.0%	99.9%	99.8%	(1.69, 1.93e-01)	<b>(8.69, 3.21e-03)</b>	<b>(7.45, 6.36e-03)</b>
15	97.0%	98.0%	99.9%	94.0%	(1.03, 3.11e-01)	<b>(13.78, 2.06e-04)</b>	<b>(5.24, 2.21e-02)</b>
21	97.0%	98.0%	97.0%	86.0%	(1.03, 3.11e-01)	(0.0, 1.)	<b>(38.89, 4.47e-10)</b>

Table 2. Comparison of the find-maximum task. The prompt simply asks GPT-4 to find the maximum of a list of numbers. Column 1: numbers uniform on (100.0, 20000.0) to 2 decimals, Column 2: numbers uniform on (100.0, 20000.0) to 12 decimals, Column 3: numbers uniform on (10, 99) as integers, Column 4: name-value pairs with values uniform on (100.0, 20000.0) to 2 decimals. The right-hand side of the table shows goodness-of-fit  $\chi^2$  tests comparing Column 1 to each of the others. Boldface entries are cases where  $p < 0.05$  and we reject the null hypothesis (that results in the given columns are produced by the same process). The hypothesis that Column 4 is produced by the same process as Column 1 is rejected: thus simply switching the list from numbers to name-value pairs introduces variance beyond what can be explained by sampling effects.

	Col. 1	Col. 2	Col. 3	Col. 4	Compare Cols(1,2) ( $\chi^2, p$ )	Compare Cols(1,3) ( $\chi^2, p$ )	Compare Cols(1,4) ( $\chi^2, p$ )
11	68.4%	70.0%	89.6%	84.8%	(0.3, 5.84e-01)	<b>(67.73, 1.88e-16)</b>	<b>(37.51, 9.08e-10)</b>
15	54.0%	48.0%	80.0%	74.0%	(3.6, 5.77e-02)	<b>(76.44, 2.27e-18)</b>	<b>(43.4, 4.46e-11)</b>
21	35.9%	41.0%	65.6%	62.7%	(2.75, 9.74e-02)	<b>(88.23, 5.83e-21)</b>	<b>(71.84, 2.34e-17)</b>

Table 3. Comparison of the find-median task. The prompt simply asks GPT-4 to find the median of a list of numbers. Column 1: numbers uniform on (100.0, 20000.0) to 2 decimals, Column 2: numbers uniform on (100.0, 20000.0) to 12 decimals, Column 3: numbers uniform on (10, 99) as integers, Column 4: name-value pairs with values uniform on (100.0, 20000.0) to 2 decimals. The right-hand side of the table shows goodness-of-fit  $\chi^2$  tests comparing Column 1 to each of the others. Boldface entries are cases where  $p < 0.05$  and we reject the null hypothesis (that results in the given columns are produced by the same process). The hypothesis that Column 3 or 4 is produced by the same process as Column 1 is rejected: thus simply changing the range on the numbers, or switching to name-value pairs introduces variance beyond what can be explained by sampling effects.

Thus, even the hypothesis that performance on the  $k_1 \times k_2$  multiplication will be equivalent to the  $k_2 \times k_1$  is rejected for at least some lengths.

Both Dziri et al (Dziri et al., 2023) and Arkoudas (Arkoudas, 2023) look at the example of long multiplication. Dziri et al note that while the answers for  $4 \times 4$  are almost always incorrect, the first and last two digits of the GPT-4 answers are almost always correct. They describe this as a matching of “surface probabilities.” That is, the first two digits of a product are determined by the leading digits of the multiplicands irrespective of length. Thus, this portion of the answer can always be determined without paying attention to the rest. Similarly for the last few digits.

### 4.5. Composite Query

Finally we examine performance on a composite query; e.g., evaluate performance on tasks that consists of sub-tasks A, B, and C and compare the resulting performance with the product of the performance on the sub-tasks. We will use the already-evaluated tasks from earlier sections as our sub-tasks. An example query is:

**Prompt:** “Find the maximum and median of this list, and then calculate their product: [24, 85, 12, 67, 55,

88, 51, 11, 66, 47, 45]. Please write ‘Answer=’”  
**Correct Answer:** 4488  
**GPT-4 Answer:** 4840

For a length  $N = 11$  list our measured value with 500 trials of this query was correct 79.0% of the time. Using GPT-4 with a list of  $N=11$  integers drawn from (10, 99) the maximum and median are correct 99.9% and 89.6% of the time respectively (these can be found from Column 3 of Tables 2 and 3 respectively). The product for 2-digit by 2-digit is correct 100.0% of the time (as shown in Table 5). Collectively, these imply that the product of maximum and median should be correct  $0.999 \times 0.896 \times 1.0$  or 89.5% of the time.

We repeat this analysis for the other lengths. For the length  $N = 15$  list our measured value was correct 69.8% of the time, versus and expected 79.9%. For the length  $N = 21$  list our measured value was correct 57.4% of the time, versus and expected 63.6%.

In each case we again have a gap between observed and expected that is larger than can be explained by a traditional estimate of margin-of-error. That performance on composite tasks is often lower than the product of the performance on sub-tasks is also observed by Dziri et al (Dziri et al., 2023).

	Col. 1	Col. 2	Col. 3	Col. 4	Compare Cols(1,2) ( $\chi^2, p$ )	Compare Cols(1,3) ( $\chi^2, p$ )	Compare Cols(1,4) ( $\chi^2, p$ )
11	98.0%	97.0%	99.7%	61.0%	(1.026, 3.11e-01)	<b>(6.356, 1.17e-02)</b>	<b>(210.002, 1.37e-47)</b>
15	91.0%	96.0%	98.0%	36.0%	<b>(10.284, 1.34e-03)</b>	<b>(23.569, 1.21e-06)</b>	<b>(326.286, 6.19e-73)</b>
21	89.0%	92.0%	99.7%	15.0%	(2.617, 1.06e-01)	<b>(53.693, 2.34e-13)</b>	<b>(548.478, 2.7e-121)</b>

Table 4. Comparison of the list-sorting task. The prompt simply asks GPT-4 to sort a list of numbers in ascending order. Column 1: numbers uniform on (100.0, 20000.0) to 2 decimals, Column 2: numbers uniform on (100.0, 20000.0) to 12 decimals, Column 3: numbers uniform on (10, 99) as integers, Column 4: name-value pairs with values uniform on (100.0, 20000.0) to 2 decimals. The right-hand side of the table shows goodness-of-fit  $\chi^2$  tests comparing Column 1 to each of the others. Boldface entries are cases where  $p < 0.05$  and we reject the null hypothesis (that results in the given columns are produced by the same process). The hypothesis that Column 3 or 4 is produced by the same process as Column 1 is rejected: thus simply changing the range on the numbers, or switching to name-value pairs introduces variance beyond what can be explained by sampling effects.

$k_1 \backslash k_2$	2	3	4	5
2	100%	90.6%	69%	40.6%
3	91.6%	55.2%	15.0%	6.2%
4	80.0%	19.4%	3.2%	1.0%
5	48.4%	8.2%	2.0%	0.0%

Table 5. Percent correct for multiplying a  $k_1$ -digit by  $k_2$ -digit number.

$k_1 \times k_2$	$k_2 \times k_1$	( $\chi^2, p$ )
$3 \times 2$	$2 \times 3$	(0.308, 0.578)
$4 \times 2$	$2 \times 4$	<b>(14.863, 1.15 e-4)</b>
$4 \times 3$	$3 \times 4$	(3.398, 0.065)
$5 \times 2$	$2 \times 5$	<b>(6.158, 0.0130)</b>
$5 \times 3$	$3 \times 5$	(1.496, 0.221)

Table 6.  $\chi^2$  goodness-of-fit test comparing the results of a  $k_1 \times k_2$  with a  $k_2 \times k_1$  multiplication (i.e., the off-diagonal elements of Table 5).

## 5. Related Work

It is well understood that the form of a prompt can greatly affect the results from a LLM as a “few-shot learner” (Brown et al., 2020), thus giving rise to the newly minted discipline of *prompt engineering*. For example, (Yu et al., 2023) show that small differences in prompting for legal reasoning tasks has a significant impact on the accuracy of responses. Our results confirm these observations for a set of simple deterministic tasks but with high statistical significance.

On the output side, Bender et al. (Bender et al., 2021) note the dangers inherent in ascribing intent and meaning to utterances generated by LLMs. In particular, we (as humans) make many assumptions about communications with other humans that can easily lead us to fall prey to the fixed-effect fallacy when working with LLMs, potentially ascribing a more general capability to the LLM than actually exists. We show that even for simple tasks there are major sources of

variance that are not easy to account for when working with LLMs.

Our experiments with deterministic algorithms are related to work that examines the capability of LLMs to perform deductive reasoning (Arkoudas, 2023). In these problems, as with most of the problems we consider, the LLM must attend to most every token in the input and not “hallucinate” new values that would lead to short-cut solutions to related but different problems than the one given. In contrast to our experiments, Arkoudas engages in a conversation with the LLM about each of the deductive problems he poses, where the LLM often proceeds to contradict itself upon getting a wrong answer. Indeed, the ad-hoc reporting of conversations with an LLM is fairly widespread (Bubeck et al., 2023) but does not rise to the level of a controlled experiment where one can make statistically significant statements. Of course, for many complex tasks it may be difficult to perform the deeper analysis we performed here for simpler tasks.

Others have observed that LLM performance degrades when the input to the LLM grows in size (within the limits of the LLM’s context window), as we have shown here. Interestingly, Liu et al (Liu et al., 2023) find that information that is at the beginning or end of the context window has more influence on LLM performance, even for simple queries that ask the LLM a question whose answer is somewhere in the input. That is, the position of information is another source of variance, as we saw in the simple prompt rewording of Table 1, where the major change was to swap the position of the input list and query (wordings 1 and 2).

Wu et al demonstrate considerable performance sensitivity for a series of tasks (Wu et al., 2023). In exploring counterfactual tasks they conclude that LLMs “rely on narrow, non-transferable procedures for task-solving.” Dziri et al explore failures of LLMs on seemingly trivial tasks (Dziri et al., 2023). They are especially interested in compositional tasks. They suggest that transformers often fail since they exploit linearized patch matching rather than any multi-step reasoning, and that errors propagate in a fashion that

compounds. Schaeffer et al suggest that the often-discussed emergent properties of LLMs are an artifact of the metrics chosen rather than any fundamental improvement (Schaeffer et al., 2023): “For a fixed task and a fixed model family, the researcher can choose a metric to create an emergent ability or choose a metric to ablate an emergent ability.”

Chain-of-Thought (CoT) is a prompting strategy that asks the LLM to output intermediate reasoning steps before giving the final answer. Research has found that it often improves LLM performance on complex tasks (Wei et al., 2022). It is worth further research to understand whether CoT-style prompts are more resilient to the variations shown in our study.

Standardized exams are often used to demonstrate LLM’s capabilities. For example, studies has shown GPT-4 achieving the passing criteria of the Japanese Medical Licensing Examination (JMLE) (Takagi et al., 2023), the Uniform Bar Examination (UBE) (Katz et al., 2023), and the US Medical Licensing Examination (USMLE) (Nori et al., 2023). Knowing that even basic tasks are sensitive to trivial variations, it is legitimate to question whether the variations between a new version of an exam and its previous versions primarily focus on factors sensitive for humans, but neglect others that can be sensitive only for LLMs.

## 6. Discussion

We’ve shown in Section 4, the risk that measured performance with a specific prompt fails to generalize to equivalent versions of the task. This work complements others that have documented the brittleness of GPT-4’s performance (see related work in Section 5). However, as far as we know, ours is the first to explore tasks with several different conditions and sufficient statistical power to rule out sampling noise as the sole source of observed variation. This allows us to state with some confidence that minor modifications have potentially enormous effects on measured capabilities. This problem is entirely orthogonal to the frequently mentioned difficulty with hallucinations.

Every measurement experiment comes with decisions about which factors might affect the output, and which should make no difference. Many of these decisions are implicit, and informed by our intuition and experience of the world. Since LLMs emulate many human capabilities it is tempting to use intuitions about humans to guide decisions about which factors should make no difference to LLM measurements. A key finding of this paper is that this assumption leads to errors that can be significant enough to invalidate claims. Bender observes that we’ve made “machines that can mindlessly generate text, but we haven’t learned how to stop imagining the mind behind it.” We suggest that the dangers of anthropomorphizing LLMs includes not just

over-interpreting their capabilities, but also imagining that their robustness to variation resembles that of humans.

An interesting direction for future work is whether we can derive new margin-of-error bounds. Our problem is that the presence of unexplained variance means that estimating  $\delta_q = 1.96 \cdot \sqrt{q \cdot (1 - q)/N}$  misses an additive component of unknown magnitude. If rewordings of a particular task can be generated automatically then estimating their variance would allow new (albeit higher) estimates of margin-of-error.

Since we warn of the risks of improper generalizations we should note the limitations of our findings. Obviously, we’ve explored a limited set of tasks, and a limited set of modifications of those tasks. The tasks in this paper are chosen deliberately with several criteria. First, they are deterministic tasks with easily-determined answers; this is clearly a very restricted portion of the problems to which LLMs might be applied. Second, the tasks we choose may be particularly difficult for transformer architectures. That is, the attention mechanism (Vaswani et al., 2017) decides which portions of the context window are most important in predicting the next token; however, for tasks like counting, sorting, etc., all words in the target list are important. Third, our prompts ask the questions in a concise and direct manner, without an attempt to guide the LLM to give a Chain-of-Thought response.

## 7. Conclusion

We have demonstrated that GPT-4 performance on simple tasks shows sensitivity to trivial modifications and that this error can be enough to invalidate claims of capabilities. Despite the limited scope of our experiments, we believe our findings point to a largely-ignored source of error that potentially affects evaluation of LLM capabilities on all tasks. That is, on every task we’ve considered we’ve found that trivial modifications introduce variance that invalidates the usual margin-of-error estimates. Our evidence doesn’t rule out the possibility that the problem might be larger, or smaller, or negligible on some other tasks. However, deciding that this source of error can be ignored for a given capability comes with a burden-of-proof, and is something that should be demonstrated empirically, rather than just assumed.

## 8. Broader impact statment

This paper presents work whose goal is to advance the field of Machine Learning. There are many potential societal consequences of our work, none which we feel must be specifically highlighted here.



## References

- Arkoudas, K. GPT-4 Can't Reason. *arXiv preprint arXiv:2308.03762*, 2023.
- Bender, E. M., Gebru, T., McMillan-Major, A., and Shmitchell, S. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pp. 610–623, 2021.
- Brown, T. B. et al. Language models are few-shot learners. In Larochelle, H., Ranzato, M., Hadsell, R., Balcan, M., and Lin, H. (eds.), *Advances in Neural Information Processing Systems 33*, 2020.
- Bubeck, S., Chandrasekaran, V., Eldan, R., Gehrke, J., Horvitz, E., Kamar, E., Lee, P., Lee, Y. T., Li, Y., Lundberg, S., et al. Sparks of artificial general intelligence: Early experiments with GPT-4. *arXiv preprint arXiv:2303.12712*, 2023.
- Clark, H. H. The language-as-fixed-effect fallacy: A critique of language statistics in psychological research. *Journal of verbal learning and verbal behavior*, 12(4):335–359, 1973.
- Coleman, E. B. Generalizing to a language population. *Psychological Reports*, 14(1):219–226, 1964.
- Dziri, N. et al. Faith and fate: Limits of transformers on compositionality. *arXiv preprint arXiv:2305.18654*, 2023.
- Gelman, A. Analysis of variance—why it is more important than ever. *Ann. Statist.*, 33(1):1–53, 2005.
- Gelman, A. and Hill, J. *Data analysis using regression and multilevel/hierarchical models*. Cambridge University Press, 2006.
- Katz, D. M., Bommarito, M. J., Gao, S., and Arredondo, P. GPT-4 passes the bar exam. *Available at SSRN 4389233*, 2023.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts. *arXiv preprint arXiv:2307.03172*, 2023.
- McCoy, R. T., Yao, S., Friedman, D., Hardy, M., and Griffiths, T. L. Embers of autoregression: Understanding large language models through the problem they are trained to solve. *arXiv preprint arXiv:2309.13638*, 2023.
- Nori, H., King, N., McKinney, S. M., Carignan, D., and Horvitz, E. Capabilities of GPT-4 on medical challenge problems. *arXiv preprint arXiv:2303.13375*, 2023.
- Schaeffer, R., Miranda, B., and Koyejo, S. Are emergent abilities of large language models a mirage?, 2023.
- Takagi, S., Watari, T., Erabi, A., Sakaguchi, K., et al. Performance of GPT-3.5 and GPT-4 on the japanese medical licensing examination: comparison study. *JMIR Medical Education*, 9(1):e48002, 2023.
- Taylor, J. R. and Thompson, W. *An introduction to error analysis: the study of uncertainties in physical measurements*, volume 2. Springer, 1982.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, Ł., and Polosukhin, I. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., Zhou, D., et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 35: 24824–24837, 2022.
- Wu, Z., Qiu, L., Ross, A., Akyürek, E., Chen, B., Wang, B., Kim, N., Andreas, J., and Kim, Y. Reasoning or reciting? exploring the capabilities and limitations of language models through counterfactual tasks, 2023.
- Yu, F., Quartey, L., and Schilder, F. Exploring the effectiveness of prompt engineering for legal reasoning tasks. In Rogers, A., Boyd-Graber, J. L., and Okazaki, N. (eds.), *Findings of the Association for Computational Linguistics: ACL 2023*, pp. 13582–13596, 2023.